



Machine Vision

CSE 489

Project 2

Submitted by:

Mohamed Abd El Hamed Moustafa - 16p8173

George Medhat Halim - 16p8197

Essam El-Din Amr - 16p9198

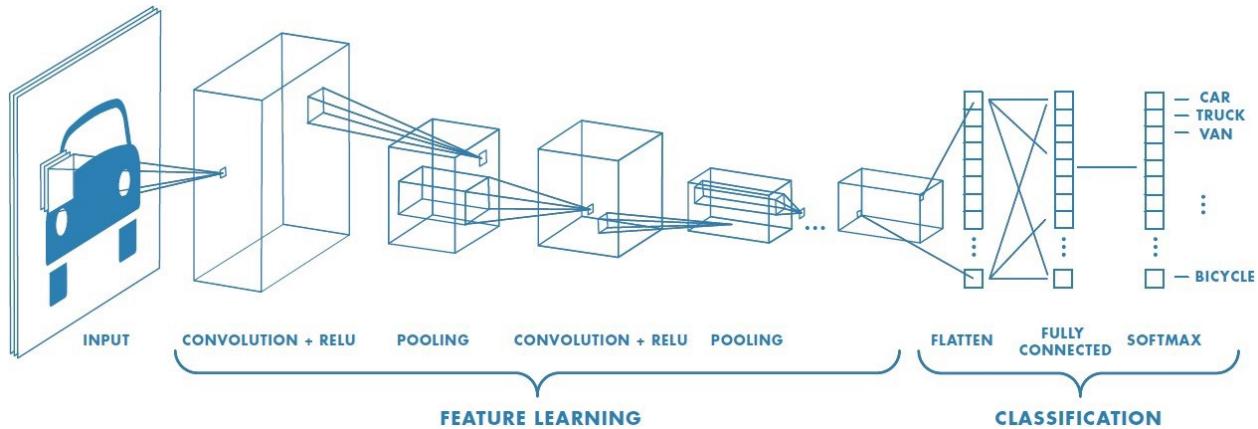
Date 12/1/2021

Senior 2

Mechatronics

Problem Definition:

The problem definition is that there are different types of objects in photos that need to be identified. The images can be of cars, cats, dogs, humans, planes ,etc. and it's required that an algorithm using Convolutional Neural Network with deep learning to classify those images and identify them with their correct label. The convolutional neural network consists of many layers usually it goes as follows, convolutional layer then max pooling layer then a flattening layer then a neural network layer then the output neural network layer. Most of those layers can be repeated many times to achieve a more complicated model to solve more complicated classification problems. Other layers can be added like dropout layers to reduce the effect of overfitting, batch normalization can also be added.



Importance:

Identifying objects using Convolutional Neural Networks s is a very important as it's one of the most well known way of classifications in the machine vision field. The classifications using Convolutional Neutral Networks has many applications like, Image recognition, Video analysis, Natural language processing, Anomaly detection, Drug discovery, Health risk assessment and biomarkers of aging discovery, Checkers game, Time series forecasting, and Cultural Heritage and 3D-datasets.

Methods and Algorithms:

Our own implemented classifier :

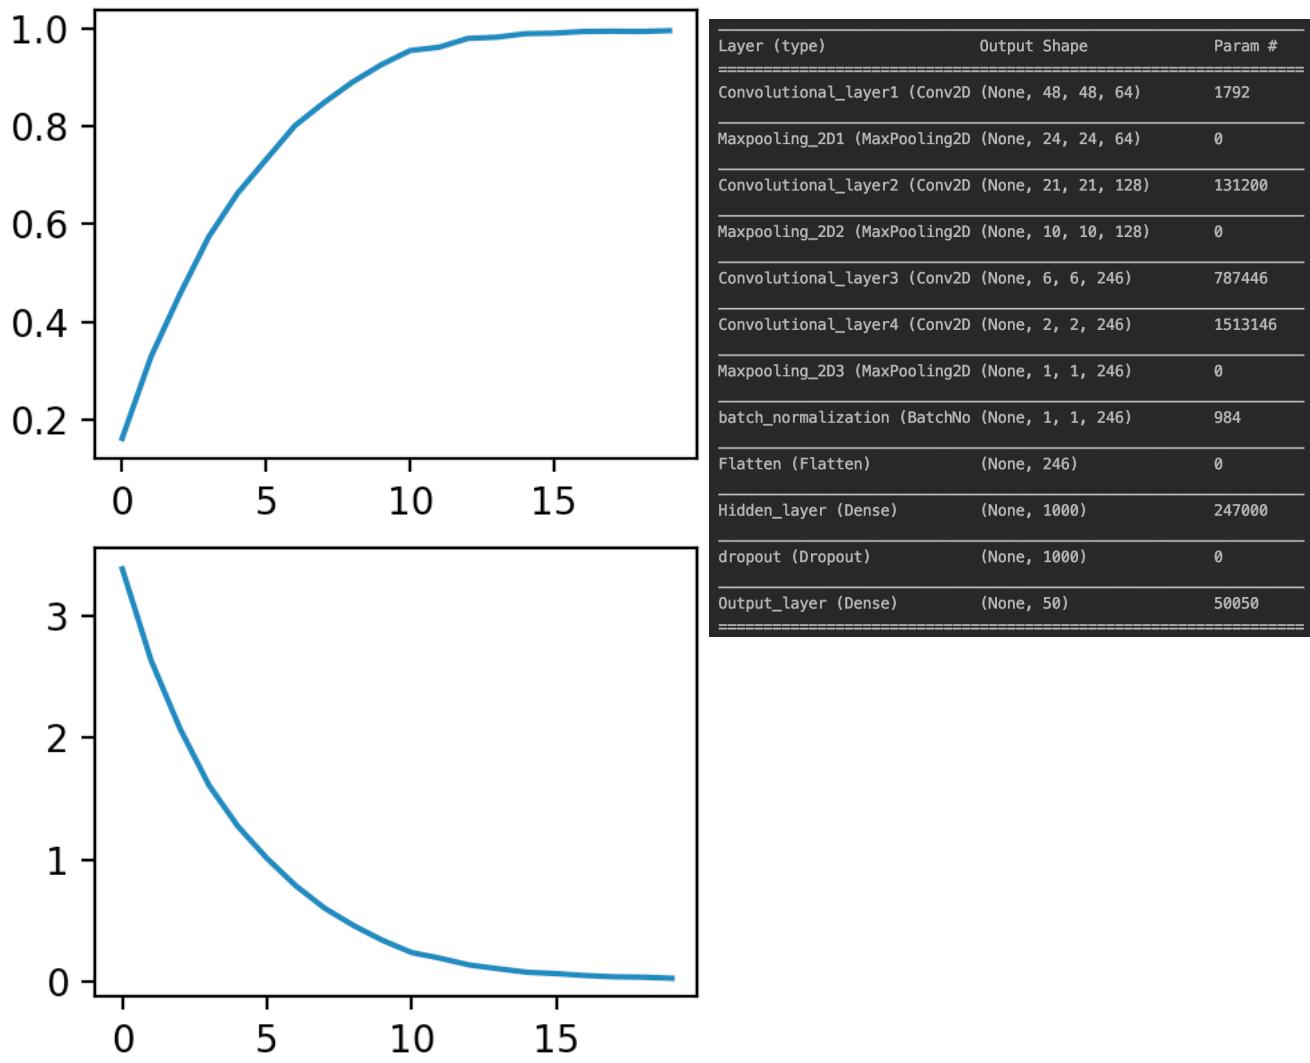
Firstly 50 class was chosen but the Catltech101 dataset has uneven images in each class so that would have trained the NN in an unbalanced way which would be unfair to some images and it would create false total accuracy as the classes with the very large training images would be learned perfectly and 15% of the total images would be taken for testing and most of them would have been classified correctly while the classes with lower number of images, the neural network would not train on them perfectly so their test accuracy would not be as good but due to their lower number it wouldn't affect the total accuracy as much so that would give false accuracy.

To over come this issue 50 images was chosen randomly for each class 36 images of them are used for training and 7 for testing and 7 for validation, 36 images for training is a very small number so mirror augmentation was used to increase the number of training images a bit. The images was also resized to be 300 by 300 pixels as not all images had the same dimensions and the has large number of pixels so the training was taking too long so it was reduced to 64 pixels by 64.

The images has been normalized and the training images has been shuffled, then a sequential model was created using 1 convolutional layer using kernel of size (3,3) and 64 filter followed by 1 max pooling layer followed by 1 convolutional layer using kernel of size (4,4) and 128 filter followed by 1 max pooling layer followed by 2 convolutional layer using kernel of size (5,5) and 246 filter followed by 1 max pooling layer followed by a batch normalization layer then a flatten layer then neural network layer of 1000 neuron followed by a drop out using 35% followed by the final output layer with 50 neurons.

The optimizer used is Adam and the loss used is the sparse categorical cross entropy and 20 epochs was used and it took about 15 minutes to train. **The Training Accuracy was 99.89% and the Validation accuracy was 50% while the testing was 52%**. Better could be achieved if i used all images in all classes as 2 or 3 classes have more than 400 image while most of them have only 50 images so when training on all images the couple 400 images classes will train very

nicely and 15 percent of those will result in great accuracy but it will falsely increase the accuracy as it's not fair as not all images trained using the same number of images and the classes that trained better have higher contribution in the total accuracy (due to have more images so 15% of those will be more than 15% of 50 images on the classes with less images) so we didn't opt to use this method and we used a fair 50 image per class as it's scientifically reasonable.



Training accuracy and loss

Testing accuracy for each class using our own implemented

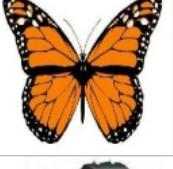
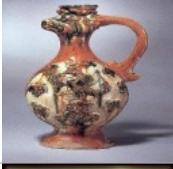
Image	Accuracy	Image	Accuracy
	57.14%		28.57%
	57.14%		42.85%
	42.85%		57.14%
	14.28%		14.28%
	100%		0%
	71.43%		42.85%
	71.43%		28.57%
	14.28%		57.14%

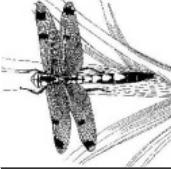
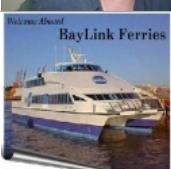
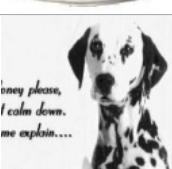
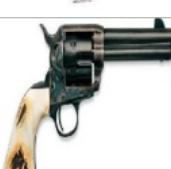
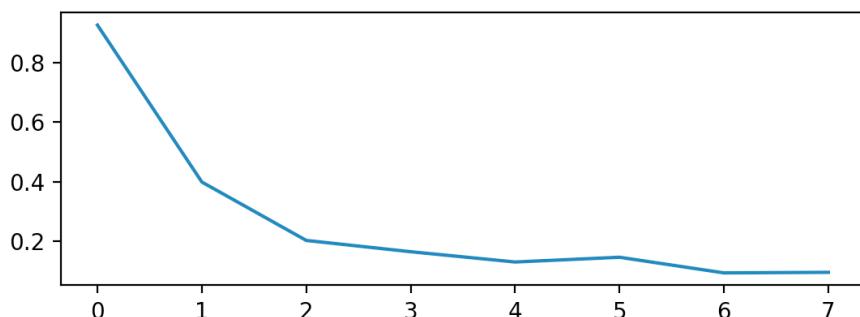
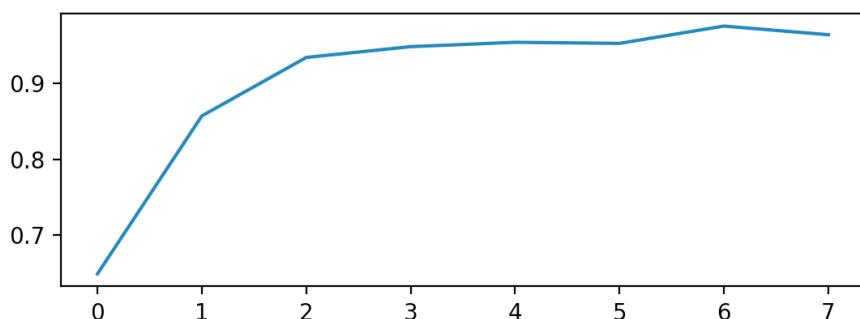
Image	Accuracy	Image	Accuracy
	71.43%		71.43%
	42.85%		0%
	28.57%		42.85%
	28.57%		28.57%
	28.57%		100%
	42.85%		85.71%
	71.42%		85.71%
	100%		28.57%
	28.57%		42.85%

Image	Accuracy	Image	Accuracy
	85.71%		57.14%
	71.42%		42.85%
	42.85%		57.14%
	57.14%		85.71%
	42.85%		85.71%
	57.14%		0%
	71.42%		100%
	42.85%		85.71%

I also tried using 200 images per category (only found 5 classes that had 200 images or more) on the model i created and the results were very promising as now it had 140 images to train on and 30 images for validation and another 30 for testing, The Training accuracy was 85% and the Testing accuracy was 82% which is very good so the short comings of the 50 class classifications is only due to the low number of images per class.

Image	Accuracy	Image	Accuracy
	100%		93.33%
	80%		96.67%
	40%		



Readymade ResNet50 classifier :

The ResNet50 model was implemented using 10 epochs. ResNet showed not very good results due to the limited number of training data and epochs as each epochs takes about 15 minutes so 10 epochs takes 2.5 hours which is very time consuming. This was done using 36 images for training ,7 for testing, and 7 for validation.

The training accuracy was 81.19% and the loss was 0.6777 while the validation accuracy was 37.43% and the loss was 3.1249 while the testing accuracy was 38.29 % and the loss was 2.9522.

Testing accuracy for each class using using ResNet50

Image	Accuracy	Image	Accuracy
	42.85%		57.14%
	57.14%		42.85%
	42.85%		28.57%
	0%		42.85%
	85.71%		42.85%

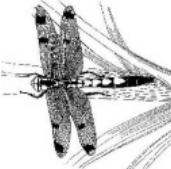
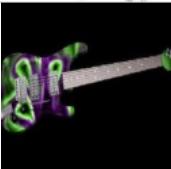
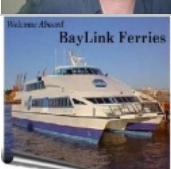
Image	Accuracy	Image	Accuracy
	14.28%		0%
	0%		0%
	28.57%		0%
	71.43%		57.14%
	14.28%		100%
	57.14%		14.28%
	57.14%		57.14%
	14.28%		28.57%
	14.28%		42.85%

Image	Accuracy	Image	Accuracy
	42.85%		0%
	0%		28.57%
	57.14%		42.85%
	42.85%		42.85%
	28.57%		71.42%
	0%		42.85%
	0%		57.14%
	14.28%		42.85%
	14.28%		85.71%

Image	Accuracy	Image	Accuracy
	100%		100%
	0%		85.71%

Readymade Pre-Trained ResNet34 classifier :

Due to the not so good results of the previous ResNet50 model, i tried another readymade model but pre trained with the ImageNet weights using transfer learning. Wide range of augmentation was used in this algorithm including and not limited to cropping, rotating, and coloring. 5 epochs was used which took almost 2.5 hours. The results were very promising the training accuracy was 99.33% and the loss was 0.0029 , and the Validation accuracy was 98.38% and the loss was 0.0035 , and the Testing accuracy was 96.887%.

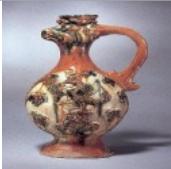
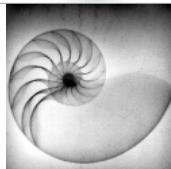
Image	Accuracy	Image	Accuracy
	100%		100%
	100%		100%
	100%		100%
	100%		100%
	100%		100%
	100%		100%

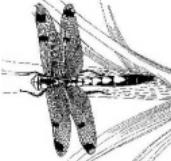
Image	Accuracy	Image	Accuracy
	100%		100%
	100%		100%
	84.21%		100%
	100%		100%
	100%		100%
	100%		100%
	100%		82.35%
	92.85%		100%
	40%		100%

Image	Accuracy	Image	Accuracy
	100%		100%
	100%		100%
	93.75%		95.83%
	89.47%		87.5%
	100%		100%
	100%		100%
	100%		100%
	100%		100%
	100%		100%

Image	Accuracy	Image	Accuracy
	100%		84.21%

Conclusion :

Our own designed CNN model showed good results given the fact that it only trained on 36 images per class while the 5 classes with 200 images each showed great results (82% on the testing images) which shows that the problem was in the low number of training data not the design itself. The ResNet50 ready-made model didn't show great results due to its complexity and time to train, the same number of epochs for our own designed CNN is about 8 times less than the time that took the ResNet50 to do the same number of epochs. While the pre-trained ResNet34 using transfers learning showed as expected great results due to its pre-trained nature using the weights from ImageNet.

Appendix:

Our own model's code:

```
1 import numpy as np
2 import cv2
3 import random
4 import matplotlib.pyplot as plt
5 from tensorflow import keras as ks
6
7 epochs = 20
8 width = 50
9 height = 50
10 histogram = np.zeros([50])
11
12 Training_images = np.zeros([1800*2, width, height, 3])
13 Training_labels = np.zeros([1800*2], dtype=int)
14 Validation_images = np.zeros([350, width, height, 3])
15 Validation_labels = np.zeros([350], dtype=int)
16 Testing_images = np.zeros([350, width, height, 3])
17 Testing_labels = np.zeros([350], dtype=int)
18
19 #index is a randomly generated list of 1800*2 number
20 #this list puts the images randomly in the training images array
21 index = list(range(0,1800*2))
22 random.shuffle(index)
23
24 indexVal = list(range(0,7*50))
25 random.shuffle(indexVal)
26
27 Train_counter = 0
28 Validation_counter = 0
29 Test_counter = 0
30 for i in range(50):
31     for j in range(50):
32         if j<36:
33             if j <9:
34                 image1 = cv2.imread("Caltech101/" + str(i + 1) + "/image_00" + str(j + 1) + ".jpg")
35                 #print("Caltech101/" + str(i + 1) + "/image_00" + str(j + 1) + ".jpg")
36             else:
37
38                 image1 = cv2.imread("Caltech101/" + str(i+1) + "/image_00" + str(j+1) + ".jpg")
39                 #print("Caltech101/" + str(i + 1) + "/image_00" + str(j + 1) + ".jpg")
40             #resize the image to be exactly 300,300 pixels as not all images had the same shape
41             image1 = image1.astype('float')
42             resized_img1 = cv2.resize(image1, (width, height), interpolation=cv2.INTER_AREA)/255
43             Training_images[[Train_counter], :, :, :] = resized_img1
44             Training_labels[[Train_counter]] = i
45             Train_counter+=Train_counter+1
46
47         if j>35 and j<43:
48             image2 = cv2.imread("Caltech101/" + str(i + 1) + "/image_00" + str(j + 1) + ".jpg")
49             image2 = image2.astype('float')
50             resized_img2 = cv2.resize(image2, (width, height), interpolation=cv2.INTER_AREA)/255
51             Validation_images[[Validation_counter], :, :, :] = resized_img2
52             Validation_labels[[Validation_counter]] = i
53             Validation_counter = Validation_counter + 1
54         if j > 42 and j < 50:
55             image3 = cv2.imread("Caltech101/" + str(i + 1) + "/image_00" + str(j + 1) + ".jpg")
56             image3 = image3.astype('float')
57             resized_img3 = cv2.resize(image3, (width, height), interpolation=cv2.INTER_AREA)/255
```

```

58             Testing_images[Test_counter,:,:,:] = resized_img3
59             Testing_labels[Test_counter] = i
60             Test_counter = Test_counter + 1
61
62         for i in range(50):
63             for j in range(36):
64
65
66             image4 = cv2.imread("Caltech101/" + str(i + 1) + "/image_000" + str(j + 51) + ".jpg")
67                 #resize the image to be exactly 300,300 pixels as not all images had the same shape
68             image4 = image4.astype('float')
69             resized_img4 = cv2.resize(image4, (width, height), interpolation=cv2.INTER_AREA)/255
70             Training_images[index[Train_counter], :, :, :] = resized_img4
71             Training_labels[index[Train_counter]] = i
72             Train_counter = Train_counter+1
73             print(Testing_labels)
74
75
76         cnn_model = ks.models.Sequential()
77         cnn_model.add(ks.layers.Conv2D(64,(3,3),activation='relu',input_shape=(width,height,3),
78                                     name='Convolutional_layer1'))
79         cnn_model.add(ks.layers.MaxPool2D((2,2),name='Maxpooling_2D1'))
80         cnn_model.add(ks.layers.Conv2D(128,(4,4),activation='relu',name='Convolutional_layer2'))
81         cnn_model.add(ks.layers.MaxPool2D((2,2),name='Maxpooling_2D2'))
82         cnn_model.add(ks.layers.Conv2D(246,(5,5),activation='relu',name='Convolutional_layer3'))
83         cnn_model.add(ks.layers.Conv2D(246,(5,5),activation='relu',name='Convolutional_layer4'))
84         cnn_model.add(ks.layers.MaxPool2D((2,2),name='Maxpooling_2D3'))
85         cnn_model.add(ks.layers.BatchNormalization())
86         cnn_model.add(ks.layers.Flatten(name='Flatten'))
87         cnn_model.add(ks.layers.Dense(1000,activation='relu',name='Hidden_layer'))
88         cnn_model.add(ks.layers.Dropout(0.35)) #0.35 was stable nausan ma
89         cnn_model.add(ks.layers.Dense(50,activation='softmax',name='Output_layer'))
90         cnn_model.summary()
91
92         cnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',metrics=['accuracy'])
93         history = cnn_model.fit(Training_images,Training_labels,epochs=epochs,
94                                 validation_data=(Validation_images,Validation_labels))
95
96
97     #Evaluate the model using the training data
98     training_loss, training_accuracy = cnn_model.evaluate(Training_images,Training_labels)
99     print('Training Accuracy {}' .format(round(float(training_accuracy),2)))
100
101    #Evaluate the model using the testing data
102    test_loss, test_accuracy = cnn_model.evaluate(Testing_images,Testing_labels)
103    print('Test Accuracy {}' .format(round(float(test_accuracy),2)))
104
105    #predictions = cnn_model.predict_classes(Testing_images)
106    predictions = np.argmax(cnn_model.predict(Testing_images), axis=-1)
107    |
108    Testing_accuracy = np.zeros([50],dtype=int)
109
110    for z in range_(350):
111        if Testing_labels[z] == predictions[z]:
112            Testing_accuracy[Testing_labels[z]] =Testing_accuracy[Testing_labels[z]] +1
113    Testing_accuracy = np.divide(Testing_accuracy,7)
114    Testing_accuracy = np.multiply(Testing_accuracy,100)
115    print(predictions)
116    print(Testing_accuracy)
117
118
119    acc = history.history['accuracy']
120    loss = history.history['loss']
121
122    fig, ax = plt.subplots(nrows=2)
123
124    ax[0].plot(acc, label="Training Accuracy")
125    ax[1].plot(loss, label="Training Loss")
126    plt.show()

```

ResNet50 Code:

```
1 import os
2 os.environ['CUDA_VISIBLE_DEVICES'] = '0'
3 import cv2
4 import numpy as np
5 from keras import layers
6 from keras.layers import Input, Add, Dense, Activation,
7 ZeroPadding2D, BatchNormalization, Flatten, Conv2D, AveragePooling2D, MaxPooling2D
8 from keras.models import Model, load_model
9 from keras.initializers import glorot_uniform
10 from keras.utils import plot_model
11 from IPython.display import SVG
12 from keras.utils.vis_utils import model_to_dot
13 import keras.backend as K
14 import tensorflow as tf
15 import random
16 from tensorflow.python.framework import ops
17
18 ROWS = 64
19 COLS = 64
20 CHANNELS = 3
21 CLASSES = 50
22
23 EPOCHS = 10
24
25 width = 64
26 height = 64
27 histogram = np.zeros([50])
28 COLAB = False
29
30 Training_images = np.zeros([1800*2,width,height,3])
31 Training_labels = np.zeros([1800*2],dtype=int)
32 Validation_images = np.zeros([350,width,height,3])
33 Validation_labels = np.zeros([350],dtype=int)
34 Testing_images = np.zeros([350,width,height,3])
35 Testing_labels = np.zeros([350],dtype=int)
36 |
37
38 #index is a randomly generated list of 1800 number
39 #this list puts the images randomly in the training images array
40 index = list(range(0,1800*2))
41 #random.seed(4)
42 random.shuffle(index)
43
44 indexVal =list(range(0,7*50))
45 random.shuffle(indexVal)
46
47 def hms_string(sec_elapsed):
48     h = int(sec_elapsed / (60 * 60))
49     m = int((sec_elapsed % (60 * 60)) / 60)
50     s = sec_elapsed % 60
51     return f"{h}:{m:02}:{s:05.2f}"
52
53 Train_counter = 0
54 Validation_counter = 0
55 Test_counter = 0
56 for i in range(50):
57     for j in range(50):
58         if j<36:
```

```

58     if j<36:
59         if j <9:
60             image1 = cv2.imread("Caltech101/" + str(i + 1) + "/image_000" +str(j + 1) + ".jpg")
61             #print("Caltech101/" + str(i + 1) + "/image_000" + str(j + 1) + ".jpg")
62         else:
63
64             image1 = cv2.imread("Caltech101/" + str(i+1) + "/image_00" + str(j+1) + ".jpg")
65             #print("Caltech101/" + str(i + 1) + "/image_00" + str(j + 1) + ".jpg")
66             #resize the image to be exactly 300,300 pixels as not all images had the same shape
67             image1 = image1.astype('float')
68             resized_img1 = cv2.resize(image1, (width,height), interpolation = cv2.INTER_AREA)/255
69             Training_images[index[Train_counter], :, :, :] = resized_img1
70             Training_labels[index[Train_counter]] = i
71             Train_counter=Train_counter+1
72
73     if j>35 and j<43:
74         image2 = cv2.imread("Caltech101/" + str(i + 1) + "/image_00" + str(j + 1) + ".jpg")
75         image2 = image2.astype('float')
76         resized_img2 = cv2.resize(image2, (width, height), interpolation=cv2.INTER_AREA)/255
77         Validation_images[indexVal[Validation_counter],:,:,:] = resized_img2
78         Validation_labels[indexVal[Validation_counter]] = i
79         Validation_counter = Validation_counter + 1
80     if j > 42 and j < 50:
81         image3 = cv2.imread("Caltech101/" + str(i + 1) + "/image_00" + str(j + 1) + ".jpg")
82         image3 = image3.astype('float')
83         resized_img3 = cv2.resize(image3, (width, height), interpolation=cv2.INTER_AREA)/255
84         Testing_images[Test_counter ,:,:,:] = resized_img3
85         Testing_labels[Test_counter ] = i
86         Test_counter = Test_counter + 1
87
88 for i in range(50):
89     for j in range(36):
90
91
92     image4 = cv2.imread("Caltech101/" + str(i + 1) + "/image_000" + str(j + 51) + ".jpg")
93         #resize the image to be exactly 300,300 pixels as not all images had the same shape
94     image4 = image4.astype('float')
95     resized_img4 = cv2.resize(image4, (width , height), interpolation = cv2.INTER_AREA)/255
96     Training_images[index[Train_counter], :, :, :] = resized_img4
97     Training_labels[index[Train_counter]] = i
98     Train_counter = Train_counter+1
99
100 def convert_to_one_hot(Y, C):
101     Y = np.eye(C)[Y.reshape(-1)].T
102     return Y
103
104 Y_train = convert_to_one_hot(Training_labels, CLASSES).T
105 Y_test = convert_to_one_hot(Testing_labels, CLASSES).T
106 Y_validation = convert_to_one_hot(Validation_labels,CLASSES).T
107
108 X_train = Training_images
109 X_test = Testing_images
110 X_validation = Validation_images

```

```

112 print ("number of training examples =", X_train.shape[0])
113 print ("number of test examples =", X_test.shape[0])
114 print ("X_train shape:", X_train.shape)
115 print ("Y_train shape:", Y_train.shape)
116 print ("X_test shape:", X_test.shape)
117 print ("Y_test shape:", Y_test.shape)
118
119 def identity_block(X, f, filters, stage, block):
120     # defining name basis
121     conv_name_base = 'res' + str(stage) + block + '_branch'
122     bn_name_base = 'bn' + str(stage) + block + '_branch'
123     F1, F2, F3 = filters
124     X_shortcut = X
125
126     X = Conv2D(filters = F1, kernel_size = (1, 1), strides = (1,1),
127                 padding = 'valid', name = conv_name_base + '2a',
128                 kernel_initializer = glorot_uniform(seed=0))(X)
129     X = BatchNormalization(axis = 3, name = bn_name_base + '2a')(X)
130     X = Activation('relu')(X)
131
132     X = Conv2D(filters = F2, kernel_size = (f, f), strides = (1,1),
133                 padding = 'same', name = conv_name_base + '2b',
134                 kernel_initializer = glorot_uniform(seed=0))(X)
135     X = BatchNormalization(axis = 3, name = bn_name_base + '2b')(X)
136     X = Activation('relu')(X)
137
138     X = Conv2D(filters = F3, kernel_size = (1, 1), strides = (1,1),
139                 padding = 'valid', name = conv_name_base + '2c',
140                 kernel_initializer = glorot_uniform(seed=0))(X)
141     X = BatchNormalization(axis = 3, name = bn_name_base + '2c')(X)
142
143     X = Add()([X, X_shortcut])
144     X = Activation('relu')(X)
145     return X
146
147 tf.compat.v1.reset_default_graph()
148
149 with tf.compat.v1.Session() as test:
150     A_prev = tf.compat.v1.placeholder("float", [3, 4, 4, 6])
151     X = np.random.randn(3, 4, 4, 6)
152     A = identity_block(A_prev, f = 2, filters = [2, 4, 6], stage = 1, block = 'a')
153     test.run(tf.compat.v1.global_variables_initializer())
154     out = test.run([A], feed_dict={A_prev: X, K.learning_phase(): 0})
155     print("out = ", out[0][1][1][0])
156
157 def convolutional_block(X, f, filters, stage, block, s=2):
158     conv_name_base = 'res' + str(stage) + block + '_branch'
159     bn_name_base = 'bn' + str(stage) + block + '_branch'
160
161     F1, F2, F3 = filters
162     X_shortcut = X
163
164     ##### MAIN PATH #####
165     X = Conv2D(F1, (1, 1), strides=(s, s), name=conv_name_base + '2a'
166     , kernel_initializer=glorot_uniform(seed=0))(X)
167     X = BatchNormalization(axis=3, name=bn_name_base + '2a')(X)
168     X = Activation('relu')(X)

```

```

169     X = Conv2D(filters=F2, kernel_size=(f, f), strides=(1, 1),
170                 padding='same', name=conv_name_base + '2b',
171                 kernel_initializer=glorot_uniform(seed=0))(X)
172     X = BatchNormalization(axis=3, name=bn_name_base + '2b')(X)
173     X = Activation('relu')(X)
174
175
176     X = Conv2D(filters=F3, kernel_size=(1, 1), strides=(1, 1),
177                 padding='valid', name=conv_name_base + '2c',
178                 kernel_initializer=glorot_uniform(seed=0))(X)
179     X = BatchNormalization(axis=3, name=bn_name_base + '2c')(X)
180
181 ###### SHORTCUT PATH #####
182 X_shortcut = Conv2D(F3, (1, 1), strides=(s, s), name=conv_name_base + '1',
183                         kernel_initializer=glorot_uniform(seed=0))(X_shortcut)
184 X_shortcut = BatchNormalization(axis=3, name=bn_name_base + '1')(X_shortcut)
185
186 # Final step: Add shortcut value to main path, and pass it through a RELU activation
187 X = Add()([X, X_shortcut])
188 X = Activation('relu')(X)
189
190 return X
191
192 tf.compat.v1.reset_default_graph()
193 with tf.compat.v1.Session() as test:
194     A_prev = tf.compat.v1.placeholder("float", [3, 4, 4, 6])
195     X = np.random.randn(3, 4, 4, 6)
196     A = convolutional_block(A_prev, f = 2, filters = [2, 4, 6], stage = 1, block = 'a')
197     test.run(tf.compat.v1.global_variables_initializer())
198     out = test.run([A], feed_dict={A_prev: X, K.learning_phase(): 0})
199     print("out = ",out[0][1][1][0])
200
201
202 def ResNet50(input_shape=(64, 64, 3), classes=2):
203     # Define the input as a tensor with shape input_shape
204     X_input = Input(input_shape)
205
206     # Zero-Padding
207     X = ZeroPadding2D((3, 3))(X_input)
208
209     # Stage 1
210     X = Conv2D(64, (7, 7), strides=(2, 2), name='conv1', kernel_initializer=glorot_uniform(seed=0))(X)
211     X = BatchNormalization(axis=3, name='bn_conv1')(X)
212     X = Activation('relu')(X)
213     X = MaxPooling2D((3, 3), strides=(2, 2))(X)
214
215     # Stage 2
216     X = convolutional_block(X, f=3, filters=[64, 64, 256], stage=2, block='a', s=1)
217     X = identity_block(X, 3, [64, 64, 256], stage=2, block='b')
218     X = identity_block(X, 3, [64, 64, 256], stage=2, block='c')
219
220     # Stage 3
221     X = convolutional_block(X, f=3, filters=[128, 128, 512], stage=3, block='a', s=2)
222     X = identity_block(X, 3, [128, 128, 512], stage=3, block='b')
223     X = identity_block(X, 3, [128, 128, 512], stage=3, block='c')
224     X = identity_block(X, 3, [128, 128, 512], stage=3, block='d')
225

```

```

184     X_shortcut = BatchNormalization(axis=3, name=bn_name_base + '1')(X_shortcut)
185
186     # Final step: Add shortcut value to main path, and pass it through a RELU activation
187     X = Add()([X, X_shortcut])
188     X = Activation('relu')(X)
189
190     return X
191
192 tf.compat.v1.reset_default_graph()
193 with tf.compat.v1.Session() as test:
194     A_prev = tf.compat.v1.placeholder("float", [3, 4, 4, 6])
195     X = np.random.randn(3, 4, 4, 6)
196     A = convolutional_block(A_prev, f = 2, filters = [2, 4, 6], stage = 1, block = 'a')
197     test.run(tf.compat.v1.global_variables_initializer())
198     out = test.run([A], feed_dict={A_prev: X, K.learning_phase(): 0})
199     print("out = ",out[0][1][1][0])
200
201
202 def ResNet50(input_shape=(64, 64, 3), classes=2):
203     # Define the input as a tensor with shape input_shape
204     X_input = Input(input_shape)
205
206     # Zero-Padding
207     X = ZeroPadding2D((3, 3))(X_input)
208
209     # Stage 1
210     X = Conv2D(64, (7, 7), strides=(2, 2), name='conv1', kernel_initializer=glorot_uniform(seed=0))(X)
211     X = BatchNormalization(axis=3, name='bn_conv1')(X)
212     X = Activation('relu')(X)
213     X = MaxPooling2D((3, 3), strides=(2, 2))(X)
214
215     # Stage 2
216     X = convolutional_block(X, f=3, filters=[64, 64, 256], stage=2, block='a', s=1)
217     X = identity_block(X, 3, [64, 64, 256], stage=2, block='b')
218     X = identity_block(X, 3, [64, 64, 256], stage=2, block='c')
219
220     # Stage 3
221     X = convolutional_block(X, f=3, filters=[128, 128, 512], stage=3, block='a', s=2)
222     X = identity_block(X, 3, [128, 128, 512], stage=3, block='b')
223     X = identity_block(X, 3, [128, 128, 512], stage=3, block='c')
224     X = identity_block(X, 3, [128, 128, 512], stage=3, block='d')
225
226     # Stage 4
227     X = convolutional_block(X, f=3, filters=[256, 256, 1024], stage=4, block='a', s=2)
228     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='b')
229     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='c')
230     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='d')
231     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='e')
232     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='f')
233
234     # Stage 5
235     X = convolutional_block(X, f=3, filters=[512, 512, 2048], stage=5, block='a', s=2)
236     X = identity_block(X, 3, [512, 512, 2048], stage=5, block='b')
237     X = identity_block(X, 3, [512, 512, 2048], stage=5, block='c')
238
239     # AVGPOOL.
240     X = AveragePooling2D((2, 2), name='avg_pool')(X)
241

```

```
241     # output layer
242     X = Flatten()(X)
243     X = Dense(classes, activation='softmax', name='fc' + str(classes), kernel_initializer=glorot_uniform(seed=0))(X)
244     # Create model
245     model = Model(inputs=X_input, outputs=X, name='ResNet50')
246
247     return model
248
249
250 model = ResNet50(input_shape = (ROWS, COLS, CHANNELS), classes = CLASSES)
251 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
252 model.fit(X_train, Y_train, epochs = EPOCHS, batch_size = 64, validation_data=(X_validation,Y_validation))
253 Testing_accuracy = np.zeros([50],dtype=int)
254 preds = model.evaluate(X_test, Y_test)
255 print(preds)
256 print(Testing_accuracy)
257 predictions = np.argmax(model.predict(Testing_images), axis=-1)
258 for z in range (350):
259     if Y_test[z] == predictions[z]:
260         Testing_accuracy[Y_test[z]] =Testing_accuracy[Y_test[z]] +1
261 Testing_accuracy = np.divide(Testing_accuracy,7)
262 Testing_accuracy = np.multiply(Testing_accuracy,100)
263
264
265 print ("Loss = " + str(preds[0]))
266 print ("Test Accuracy = " + str(preds[1]))
267 model.summary()
268 |
```

ResNet34 Pre-Trained Code:

```
1  import matplotlib.pyplot as plt
2  import matplotlib
3  import cv2
4  import os
5  import torch
6  import numpy as np
7  import torch.nn as nn
8  import torch.nn.functional as F
9  import torch.optim as optim
10 import time
11 import random
12 import pretrainedmodels
13 from imutils import paths
14 from sklearn.preprocessing import LabelBinarizer
15 from sklearn.model_selection import train_test_split
16 from torchvision.transforms import transforms
17 from torch.utils.data import DataLoader, Dataset
18 from tqdm import tqdm
19
20 matplotlib.style.use('ggplot')
21 '''SEED Everything'''
22 |
23 def seed_everything(SEED=42):
24     random.seed(SEED)
25     np.random.seed(SEED)
26     torch.manual_seed(SEED)
27     torch.cuda.manual_seed(SEED)
28     torch.cuda.manual_seed_all(SEED)
29     torch.backends.cudnn.benchmark = True # keep True if all the input have same size.
30
```

```
32 SEED = 42
33 seed_everything(SEED)
34 '''SEED Everything'''
35
36 if torch.cuda.is_available():
37     device = 'cuda'
38 else:
39     device = 'cpu'
40
41 epochs = 5
42 BATCH_SIZE = 16
43
44 image_paths = list(paths.list_images('101_ObjectCategories'))
45 data = []
46 labels = []
47 label_names = []
48 for image_path in image_paths:
49     label = image_path.split(os.path.sep)[-2]
50     if label == 'BACKGROUND_Google':
51         continue
52     image = cv2.imread(image_path)
53     image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
54     data.append(image)
55     label_names.append(label)
56     labels.append(label)
57 data = np.array(data)
58 labels = np.array(labels)
59 print(f"Labels: {labels[0]}")
60
61 lb = LabelBinarizer()
62 labels = lb.fit_transform(labels)
63 print(f"Total number of classes: {len(lb.classes_)}")
```

```

65     count_arr = []
66     label_arr = []
67     for i in range(len(lb.classes_)):
68         count = 0
69         # print(lb.classes_[i])
70         for j in range(len(label_names)):
71             if lb.classes_[i] in label_names[j]:
72                 count += 1
73         count_arr.append(count)
74     label_arr.append(lb.classes_[i])
75
76     plt.figure(figsize=(18, 15))
77     plt.bar(label_arr, count_arr, )
78     plt.xticks(rotation='vertical')
79     #plt.show()
80
81     (X, x_val, Y, y_val) = train_test_split(data, labels,
82                                              test_size=0.16,
83                                              stratify=labels,
84                                              random_state=42)
85     (x_train, x_test, y_train, y_test) = train_test_split(X, Y,
86                                              test_size=0.1903,
87                                              random_state=42)
88     print(f"x_train examples: {x_train.shape}\nx_test examples: {x_test.shape}\nx_val examples: {x_val.shape}")
89
90     train_transform = transforms.Compose(
91         [transforms.ToPILImage(),
92          transforms.Resize((224, 224)),
93          transforms.ToTensor(),
94          transforms.Normalize(mean=[0.485, 0.456, 0.406],
95                               std=[0.229, 0.224, 0.225])))
96     val_transform = transforms.Compose(
97         [transforms.ToPILImage(),
98          transforms.Resize((224, 224)),
99          transforms.ToTensor(),
100         transforms.Normalize(mean=[0.485, 0.456, 0.406],
101                           std=[0.229, 0.224, 0.225])))
102
103
104     # custom dataset
105     class ImageDataset(Dataset):
106         def __init__(self, images, labels=None, transforms=None):
107             self.X = images
108             self.y = labels
109             self.transforms = transforms
110
111         def __len__(self):
112             return (len(self.X))
113
114         def __getitem__(self, i):
115             data = self.X[i][:]
116
117             if self.transforms:
118                 data = self.transforms(data)
119
120             if self.y is not None:
121                 return (data, self.y[i])
122             else:
123                 return data
124
125
126     train_data = ImageDataset(x_train, y_train, train_transform)
127     val_data = ImageDataset(x_val, y_val, val_transform)
128     test_data = ImageDataset(x_test, y_test, val_transform)

```

```

130     # dataloaders
131     trainloader = DataLoader(train_data, batch_size=16, shuffle=True)
132     valloader = DataLoader(val_data, batch_size=16, shuffle=True)
133     testloader = DataLoader(test_data, batch_size=16, shuffle=False)
134
135
136     class ResNet34(nn.Module):
137         def __init__(self, pretrained):
138             super(ResNet34, self).__init__()
139             if pretrained is True:
140                 self.model = pretrainedmodels.__dict__['resnet34'](pretrained='imagenet')
141             else:
142                 self.model = pretrainedmodels.__dict__['resnet34'](pretrained=None)
143
144             self.l0 = nn.Linear(512, len(lb.classes_))
145             self.dropout = nn.Dropout2d(0.4)
146
147         def forward(self, x):
148             # get the batch size only, ignore (c, h, w)
149             batch, _, _, _ = x.shape
150             x = self.model.features(x)
151             x = F.adaptive_avg_pool2d(x, 1).reshape(batch, -1)
152             x = self.dropout(x)
153             l0 = self.l0(x)
154             return l0
155
156
157     model = ResNet34(pretrained=True).to(device)
158     print(model)
159
160     # optimizer
161     optimizer = optim.Adam(model.parameters(), lr=1e-4)
162     # loss function

```

```

163     criterion = nn.CrossEntropyLoss()
164
165     # validation function
166     def validate(model, dataloader):
167         print('Validating')
168         model.eval()
169         val_running_loss = 0.0
170         val_running_correct = 0
171         with torch.no_grad():
172             for i, data in tqdm(enumerate(dataloader), total=int(len(val_data) / dataloader.batch_size)):
173                 data, target = data[0].to(device), data[1].to(device)
174                 outputs = model(data)
175                 loss = criterion(outputs, torch.max(target, 1)[1])
176
177                 val_running_loss += loss.item()
178                 _, preds = torch.max(outputs.data, 1)
179                 val_running_correct += (preds == torch.max(target, 1)[1]).sum().item()
180
181             val_loss = val_running_loss / len(dataloader.dataset)
182             val_accuracy = 100. * val_running_correct / len(dataloader.dataset)
183             print(f'Val Loss: {val_loss:.4f}, Val Acc: {val_accuracy:.2f}')
184
185         return val_loss, val_accuracy
186
187     # training function
188     def fit(model, dataloader):
189         print('Training')
190         model.train()
191         train_running_loss = 0.0
192         train_running_correct = 0
193         for i, data in tqdm(enumerate(dataloader), total=int(len(train_data) / dataloader.batch_size)):
194             data, target = data[0].to(device), data[1].to(device)

```

```

195     optimizer.zero_grad()
196     outputs = model(data)
197     # print(outputs)
198     # print(torch.max(target, 1)[1])
199     loss = criterion(outputs, torch.max(target, 1)[1])
200     train_running_loss += loss.item()
201     _, preds = torch.max(outputs.data, 1)
202     train_running_correct += (preds == torch.max(target, 1)[1]).sum().item()
203     loss.backward()
204     optimizer.step()
205
206     train_loss = train_running_loss / len(dataloader.dataset)
207     train_accuracy = 100. * train_running_correct / len(dataloader.dataset)
208
209     print(f"Train Loss: {train_loss:.4f}, Train Acc: {train_accuracy:.2f}")
210
211     return train_loss, train_accuracy
212
213
214 def class_acc_test(model, testloader):
215     class_correct = list(0. for _ in range(len(lb.classes_)))
216     class_total = list(0. for _ in range(len(lb.classes_)))
217     with torch.no_grad():
218         for data in testloader:
219             inputs, target = data[0].to(device), data[1].to(device)
220             outputs = model(inputs)
221             _, predicted = torch.max(outputs.data, 1)
222             correct = (predicted == torch.max(target, 1)[1])
223             for i in range(len(predicted)):
224                 label = torch.max(target, 1)[1][i]
225                 class_correct[label] += correct[i].item()
226                 class_total[label] += 1
227             for i in range(len(lb.classes_)):
228                 print(f"Accuracy of {lb.classes_[i]}: {100 * (class_correct[i] / class_total[i]):.3f}")
229
230
231     train_loss, train_accuracy = [], []
232     val_loss, val_accuracy = [], []
233     print(f"Training on {len(train_data)} examples, validating on {len(val_data)} examples...")
234     start = time.time()
235     for epoch in range(epochs):
236         print(f"Epoch {epoch + 1} of {epochs}")
237         train_epoch_loss, train_epoch_accuracy = fit(model, trainloader)
238         val_epoch_loss, val_epoch_accuracy = validate(model, valloader)
239         train_loss.append(train_epoch_loss)
240         train_accuracy.append(train_epoch_accuracy)
241         val_loss.append(val_epoch_loss)
242         val_accuracy.append(val_epoch_accuracy)
243     end = time.time()
244
245     print((end - start) / 60, 'minutes')
246
247     # torch.save(model.state_dict(), f"../outputs/models/resnet34_{epochs}.pth")
248     # accuracy plots
249     plt.figure(figsize=(10, 7))
250     plt.plot(train_accuracy, color='green', label='train accuracy')
251     plt.plot(val_accuracy, color='blue', label='validation accuracy')
252     plt.xlabel('Epochs')
253     plt.ylabel('Accuracy')
254     plt.legend()
255     plt.savefig('Accuracy.png')
256     plt.show()

```

```
258 # loss plots
259 plt.figure(figsize=(10, 7))
260 plt.plot(train_loss, color='orange', label='train loss')
261 plt.plot(val_loss, color='red', label='validation loss')
262 plt.xlabel('Epochs')
263 plt.ylabel('Loss')
264 plt.legend()
265 plt.savefig('loss.png')
266 plt.show()
267
268 def test(model, dataloader):
269     correct = 0
270     total = 0
271     with torch.no_grad():
272         for data in testloader:
273             inputs, target = data[0].to(device), data[1].to(device)
274             outputs = model(inputs)
275             _, predicted = torch.max(outputs.data, 1)
276             total += target.size(0)
277             correct += (predicted == torch.max(target, 1)[1]).sum().item()
278     return correct, total
279
280
281 correct, total = test(model, testloader)
282 print('Accuracy of the network on test images: %0.3f %%' % (100 * correct / total))
283 class_acc_test(model, testloader)
284
```